

# Advanced DevOps Lab Experiment 4

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy  
Your First Kubernetes Application.

Theory:

Kubernetes, originally developed by Google, is an open-source container orchestration platform. It automates the deployment, scaling, and management of containerized applications, ensuring high availability and fault tolerance. Kubernetes is now the industry standard for container orchestration and is governed by the Cloud Native Computing Foundation (CNCF), with contributions from major cloud and software providers like Google, AWS, Microsoft, IBM, Intel, Cisco, and Red Hat.

Kubernetes Deployment: Is a resource in Kubernetes that provides declarative updates for Pods and ReplicaSets. With a Deployment, you can define how many replicas of a pod should run, roll out new versions of an application, and roll back to previous versions if necessary. It ensures that the desired number of pod replicas are running at all times.

Necessary Requirements:

- EC2 Instance: The experiment required launching a t2.medium EC2 instance with 2 CPUs, as Kubernetes demands sufficient resources for effective functioning.
- Minimum Requirements:
  - Instance Type: t2.medium
  - CPUs: 2
  - Memory: Adequate for container orchestration.

This ensured that the Kubernetes cluster had the necessary resources to function smoothly.

Note:

AWS Personal Account is preferred but we can also perform it on AWS Academy(adding some ignores in the command if any error occurs in below as the below experiment is performed on Personal Account  
.).

If You are using AWS Academy Account Errors you will face in kubeadm init command so you have to

add some ignores with this command.

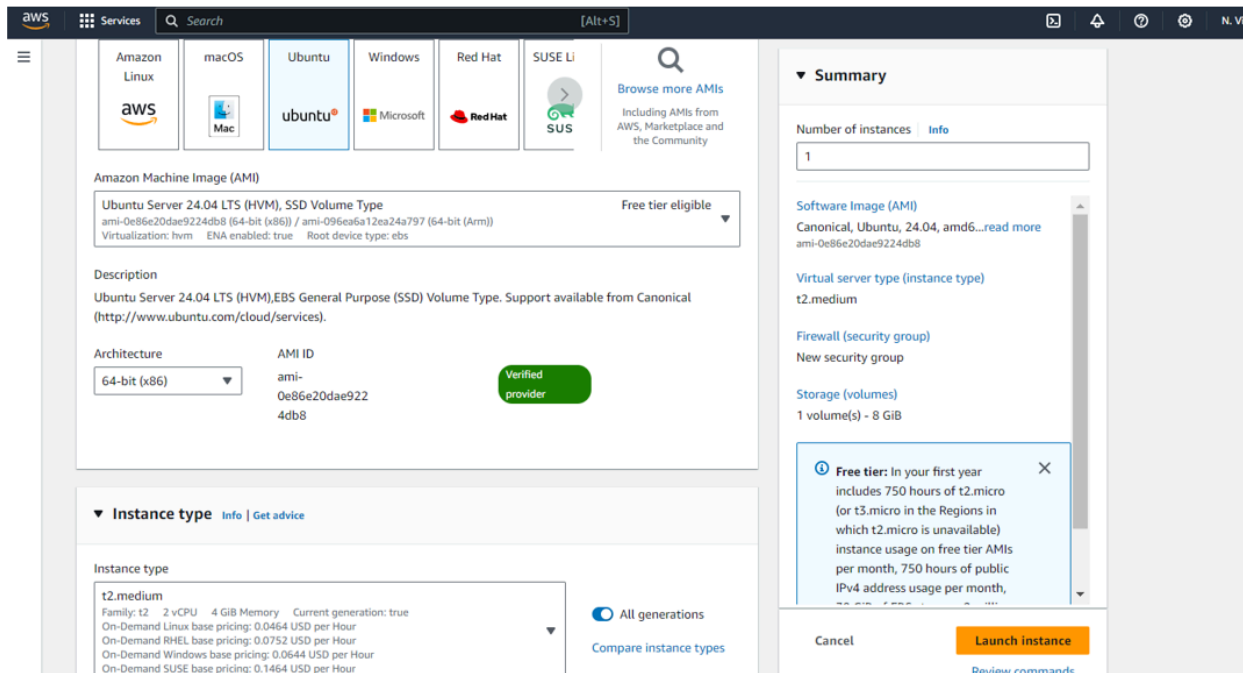
Step 1: Log in to your AWS Academy/personal account and launch a new Ec2 Instance.

Select Ubuntu as AMI and t2.medium as Instance Type, create a key of type RSA with .pem extension,

and move the downloaded key to the new folder.

Note: A minimum of 2 CPUs are required so Please select t2.medium and do not forget to stop the

instance after the experiment because it is not available in the free tier.



Step 2: After creating the instance click on Connect the instance and navigate to SSH Client.

Instances (1) Info Last updated less than a minute ago Refresh Connect Instance state Actions Launch instances

Find Instance by attribute or tag (case-sensitive) All states

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IP
Exp4	i-04a49a928f6d2dbf1	Running	t2.medium	Initializing	View alarms	us-east-1d	ec2-52-...

---

aws Services Search [Alt+S]

## Connect to instance Info

Connect to your instance i-04a49a928f6d2dbf1 (Exp4) using any of these options

EC2 Instance Connect | Session Manager | **SSH client** | EC2 serial console

Instance ID  
i-04a49a928f6d2dbf1 (Exp4)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is exp4key.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.  
`chmod 400 "exp4key.pem"`
4. Connect to your instance using its Public DNS:  
`ec2-52-91-240-34.compute-1.amazonaws.com`

Example:  
`ssh -i "exp4key.pem" ubuntu@ec2-52-91-240-34.compute-1.amazonaws.com`

**Note:** In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Step 3: Now open the folder in the terminal where our .pem key is stored and paste the Example

command (starting with ssh -i ..... ) in the terminal.( ssh -i "Master\_Ec2\_Key.pem" ubuntu@ec2-54-196-129-215.compute-1.amazonaws.com)

```
Microsoft Windows [Version 10.0.22000.2057]
(c) Microsoft Corporation. All rights reserved.

C:\Users\ACER\Downloads>ssh -i "exp4key.pem" ubuntu@ec2-52-91-240-34.compute-1.amazonaws.com
The authenticity of host 'ec2-52-91-240-34.compute-1.amazonaws.com (52.91.240.34)' can't be established.
ECDSA key fingerprint is SHA256:hQXGxhM3JrUApDQWob0ui+rTZu/uzA7hY4Hs9p58oLM.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-52-91-240-34.compute-1.amazonaws.com,52.91.240.34' (ECDSA) to the list of known hosts.
Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/pro

System information as of Sat Sep 21 10:54:46 UTC 2024

System load:  0.08               Processes:    115
Usage of /:   22.7% of 6.71GB    Users logged in: 0
Memory usage: 6%                IPv4 address for enX0: 172.31.87.78
Swap usage:  0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status
```

Step 4: Run the below commands to install and setup Docker.

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee
/etc/apt/trusted.gpg.d/docker.gpg > /dev/null
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu
$(lsb_release -cs) stable"
sudo apt-get update
sudo apt-get install -y docker-ce
```

```
9AcZ58Em+1WsVnAXdUR//bMmhyr8wL/G1YO1V3JEJTRdxsSxdYa4deGBBY/Adpsw
24jxhO3JR+1sJpqIUeb999+R8euDhRHG9eF07DRu6weatUJ6suupoDTRWtr/4yGqe
dKxV3qQhNLsNaAzqW/1nA3iUB4k7KCaKZxhdhDbClf9P37qaRW467BLCV0/coL3y
Vm50dwdrNTKpMBh3ZpbB1uJvgi9mXtyBOMJ3v8RZeDzFiG8HdCtg9RvIt/AIFoHR
H3S+U79NT6i0KPz1ImDfs8T7R1pyuMc4Ufs8ggyg9v3Ae6cN3eQyxcK3w0cbBwsh
/nQNFsA6uu+9H7NhbehBMhYnpNZyrHzCmzyXkaumRAQoCbGCNykTRwsur9gS41TQ
M8ssD1jFheO3f3h0DnkKU+HKjvMR01DK7zdmLdNZ1cvtZH/nCC9KPj1z8QC47S
xx+dTZ5x40NAhw5/LN3PoKtn8LPjY9NP9uDWI+TwYquS2U+KHDrBD1sgozDbs/0
jCxcPdzNmXpWQHETHU76490XHP7UeNST1mCUC5qdank0V1ieJf6/CfTFU4MfcrG
YT90qFF93M3v01BbxP+EIY2/9tiIPbrd
=0YYh
-----END PGP PUBLIC KEY BLOCK-----
ubuntu@ip-172-31-87-78:~$ /etc/apt/trusted.gpg.d/docker.gpg > /dev/null
-bash: /etc/apt/trusted.gpg.d/docker.gpg: No such file or directory
ubuntu@ip-172-31-87-78:~$ sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
Repository: 'deb [arch=amd64] https://download.docker.com/linux/ubuntu noble stable'
Description:
Archive for codename: noble components: stable
More info: https://download.docker.com/linux/ubuntu
Adding repository.
Press [ENTER] to continue or Ctrl-c to cancel.
Adding deb entry to /etc/apt/sources.list.d/archive_uri-https_download_docker_com_linux_ubuntu-noble.list
Adding disabled deb-src entry to /etc/apt/sources.list.d/archive_uri-https_download_docker_com_linux_ubuntu-noble.list
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease [126 kB]
Get:4 http://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]
Get:5 https://download.docker.com/linux/ubuntu noble InRelease [48.8 kB]
Get:6 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 Packages [15.0 MB]
Get:7 https://download.docker.com/linux/ubuntu noble/stable amd64 Packages [15.3 kB]
Get:8 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe Translation-en [5982 kB]
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details
ubuntu@ip-172-31-87-78:~$ sudo apt-get update
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Hit:5 http://security.ubuntu.com/ubuntu noble-security InRelease
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details
ubuntu@ip-172-31-87-78:~$ sudo apt-get install -y docker-ce
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  containerd.io docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libbtdl7 libslirp0 pigz slirp4netns
Suggested packages:
  aufs-tools cgroupfs-mount | cgroup-lite
The following NEW packages will be installed:
  containerd.io docker-buildx-plugin docker-ce docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libbtdl7 libslirp0 pigz slirp4netns
0 upgraded, 10 newly installed, 0 to remove and 139 not upgraded.
Need to get 123 MB of archives.
After this operation, 442 MB of additional disk space will be used.
Get:1 https://download.docker.com/linux/ubuntu noble/stable amd64 containerd.io 1.7.16~3-ubuntu~20.04~20230801 [17.5 MB]
Get:2 https://download.docker.com/linux/ubuntu noble/stable amd64 docker-buildx-plugin 0.11.2~3-ubuntu~20.04~20230801 [2.8 MB]
Get:3 https://download.docker.com/linux/ubuntu noble/stable amd64 docker-ce-cli 26.1.1~3-ubuntu~20.04~20230801 [4.5 MB]
Get:4 https://download.docker.com/linux/ubuntu noble/stable amd64 docker-ce-rootless-extras 26.1.1~3-ubuntu~20.04~20230801 [1.2 MB]
Get:5 https://download.docker.com/linux/ubuntu noble/stable amd64 docker-compose-plugin 2.20.1~3-ubuntu~20.04~20230801 [4.5 MB]
Get:6 https://download.docker.com/linux/ubuntu noble/stable amd64 libbtdl7 1.4.4~3-ubuntu~20.04~20230801 [1.2 MB]
Get:7 https://download.docker.com/linux/ubuntu noble/stable amd64 libslirp0 4.6.0~3-ubuntu~20.04~20230801 [1.2 MB]
Get:8 https://download.docker.com/linux/ubuntu noble/stable amd64 pigz 2.17-1 [1.2 MB]
Get:9 https://download.docker.com/linux/ubuntu noble/stable amd64 slirp4netns 1.1.8-3 [1.2 MB]
Get:10 https://download.docker.com/linux/ubuntu noble/stable amd64 docker-ce 26.1.1~3-ubuntu~20.04~20230801 [1.2 MB]
debconf: delaying package configuration, since apt-utils is not installed
Fetched 123 MB of archives
Unpacking containerd.io (1.7.16~3-ubuntu~20.04~20230801) ...
Unpacking docker-buildx-plugin (0.11.2~3-ubuntu~20.04~20230801) ...
Unpacking docker-ce-cli (26.1.1~3-ubuntu~20.04~20230801) ...
Unpacking docker-ce-rootless-extras (26.1.1~3-ubuntu~20.04~20230801) ...
Unpacking docker-compose-plugin (2.20.1~3-ubuntu~20.04~20230801) ...
Unpacking libbtdl7 (1.4.4~3-ubuntu~20.04~20230801) ...
Unpacking libslirp0 (4.6.0~3-ubuntu~20.04~20230801) ...
Unpacking pigz (2.17-1) ...
Unpacking slirp4netns (1.1.8-3) ...
Unpacking docker-ce (26.1.1~3-ubuntu~20.04~20230801) ...
Setting up containerd.io (1.7.16~3-ubuntu~20.04~20230801) ...
Setting up docker-buildx-plugin (0.11.2~3-ubuntu~20.04~20230801) ...
Setting up docker-ce-cli (26.1.1~3-ubuntu~20.04~20230801) ...
Setting up docker-ce-rootless-extras (26.1.1~3-ubuntu~20.04~20230801) ...
Setting up docker-compose-plugin (2.20.1~3-ubuntu~20.04~20230801) ...
Setting up libbtdl7 (1.4.4~3-ubuntu~20.04~20230801) ...
Setting up libslirp0 (4.6.0~3-ubuntu~20.04~20230801) ...
Setting up pigz (2.17-1) ...
Setting up slirp4netns (1.1.8-3) ...
Setting up docker-ce (26.1.1~3-ubuntu~20.04~20230801) ...
```

```

Unpacking docker-ce (5:27.3.1-1~ubuntu.24.04~noble) ...
Selecting previously unselected package docker-ce-rootless-extras.
Preparing to unpack .../5-docker-ce-rootless-extras_5%3a27.3.1-1~ubuntu.24.04~noble_amd64.deb ...
Unpacking docker-ce-rootless-extras (5:27.3.1-1~ubuntu.24.04~noble) ...
Selecting previously unselected package docker-compose-plugin.
Preparing to unpack .../6-docker-compose-plugin_2.29.7-1~ubuntu.24.04~noble_amd64.deb ...
Unpacking docker-compose-plugin (2.29.7-1~ubuntu.24.04~noble) ...
Selecting previously unselected package libltdl7:amd64.
Preparing to unpack .../7-libltdl7_2.4.7-7build1_amd64.deb ...
Unpacking libltdl7:amd64 (2.4.7-7build1) ...
Selecting previously unselected package libslirp0:amd64.
Preparing to unpack .../8-libslirp0_4.7.0-1ubuntu3_amd64.deb ...
Unpacking libslirp0:amd64 (4.7.0-1ubuntu3) ...
Selecting previously unselected package slirp4netns.
Preparing to unpack .../9-slirp4netns_1.2.1-1build2_amd64.deb ...
Unpacking slirp4netns (1.2.1-1build2) ...
Setting up docker-buildx-plugin (0.17.1-1~ubuntu.24.04~noble) ...
Setting up containerd.io (1.7.22-1) ...
Created symlink /etc/systemd/system/multi-user.target.wants/containerd.service → /usr/lib/systemd/system/containerd.service.
Setting up docker-compose-plugin (2.29.7-1~ubuntu.24.04~noble) ...
Setting up libltdl7:amd64 (2.4.7-7build1) ...
Setting up docker-ce-cli (5:27.3.1-1~ubuntu.24.04~noble) ...
Setting up libslirp0:amd64 (4.7.0-1ubuntu3) ...
Setting up pigz (2.8-1) ...
Setting up docker-ce-rootless-extras (5:27.3.1-1~ubuntu.24.04~noble) ...
Setting up slirp4netns (1.2.1-1build2) ...
Setting up docker-ce (5:27.3.1-1~ubuntu.24.04~noble) ...
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.
Processing triggers for man-db (2.12.0-4build2) ...
Processing triggers for libc-bin (2.39-0ubuntu8.2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.

```

`sudo mkdir -p /etc/docker`

`cat <<EOF | sudo tee /etc/docker/daemon.json`

```

{
"exec-opts": ["native.cgroupdriver=systemd"]
}
EOF

```

```

ubuntu@ip-172-31-87-78:~$ sudo mkdir -p /etc/docker
ubuntu@ip-172-31-87-78:~$ cat <<EOF | sudo tee /etc/docker/daemon.json
> {
> "exec-opts": ["native.cgroupdriver=systemd"]
> }
> EOF
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
ubuntu@ip-172-31-87-78:~$ █

```

`sudo systemctl enable docker`

`sudo systemctl daemon-reload`

`sudo systemctl restart docker`

```
ubuntu@ip-172-31-87-78:~$ sudo systemctl enable docker
Synchronizing state of docker.service with SysV service script with /usr/lib/systemd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable docker
ubuntu@ip-172-31-87-78:~$ sudo systemctl daemon-reload
ubuntu@ip-172-31-87-78:~$ sudo systemctl restart docker
```

```
ubuntu@ip-172-31-87-78:~$ sudo systemctl restart docker
ubuntu@ip-172-31-87-78:~$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
ubuntu@ip-172-31-87-78:~$ echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /
ubuntu@ip-172-31-87-78:~$
```

Step 5: Run the below command to install Kubernetes.

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
sudo apt-get update
sudo apt-get install -y kubelet kubeadm kubectl
sudo apt-mark hold kubelet kubeadm kubectl
```

```
ubuntu@ip-172-31-87-78:~$ sudo apt-get update
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Get:5 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb InRelease [1186 B]
Hit:6 http://security.ubuntu.com/ubuntu noble-security InRelease
Get:7 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb Packages [4865 B]
Fetched 6051 B in 1s (10.5 kB/s)
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), s
ubuntu@ip-172-31-87-78:~$ sudo apt-get install -y kubelet kubeadm kubectl
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  conntrack cri-tools kubernetes-cni
The following NEW packages will be installed:
  conntrack cri-tools kubeadm kubectl kubelet kubernetes-cni
0 upgraded, 6 newly installed, 0 to remove and 139 not upgraded.
Need to get 87.4 MB of archives.
After this operation, 314 MB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 conntrack amd64 1:1.4.8-1ubuntu1 [37.9 kB]
Get:2 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb cri-tools 1.31.1-1.1 [15.7 MB]
Get:3 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb kubeadm 1.31.1-1.1 [11.4 MB]
Get:4 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb kubectl 1.31.1-1.1 [11.2 MB]
Get:5 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb kubernetes-cni 1.5.1-1.1 [33.9 MB]
Get:6 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb kubelet 1.31.1-1.1 [15.2 MB]
Fetched 87.4 MB in 1s (87.2 MB/s)
Selecting previously unselected package conntrack.
(Reading database ... 68007 files and directories currently installed.)
Preparing to unpack .../0-conntrack_1%3a1.4.8-1ubuntu1_amd64.deb ...
Unpacking conntrack (1:1.4.8-1ubuntu1) ...
Selecting previously unselected package cri-tools.
Preparing to unpack .../1-cri-tools_1.31.1-1.1_amd64.deb ...
Unpacking cri-tools (1.31.1-1.1) ...
Selecting previously unselected package kubeadm.
Preparing to unpack .../2-kubeadm_1.31.1-1.1_amd64.deb ...
Unpacking kubeadm (1.31.1-1.1) ...
Selecting previously unselected package kubectl.
Preparing to unpack .../3-kubectl_1.31.1-1.1_amd64.deb ...
Unpacking kubectl (1.31.1-1.1) ...
Selecting previously unselected package kubernetes-cni.
Preparing to unpack .../4-kubernetes-cni_1.5.1-1.1_amd64.deb ...
Unpacking kubernetes-cni (1.5.1-1.1) ...
Selecting previously unselected package kubelet.
Preparing to unpack .../5-kubelet_1.31.1-1.1_amd64.deb ...
Unpacking kubelet (1.31.1-1.1) ...
```

```

get:b https://prod-cdn.packages.k8s.io/repositories/1sv:/kubernetes:/core:/stable:/v1.31/deb kubelet 1.
Fetched 87.4 MB in 1s (87.2 MB/s)
Selecting previously unselected package conntrack.
(Reading database ... 68007 files and directories currently installed.)
Preparing to unpack .../0-conntrack_1%3a1.4.8-1ubuntu1_amd64.deb ...
Unpacking conntrack (1:1.4.8-1ubuntu1) ...
Selecting previously unselected package cri-tools.
Preparing to unpack .../1-cri-tools_1.31.1-1.1_amd64.deb ...
Unpacking cri-tools (1.31.1-1.1) ...
Selecting previously unselected package kubeadm.
Preparing to unpack .../2-kubeadm_1.31.1-1.1_amd64.deb ...
Unpacking kubeadm (1.31.1-1.1) ...
Selecting previously unselected package kubect1.
Preparing to unpack .../3-kubect1_1.31.1-1.1_amd64.deb ...
Unpacking kubect1 (1.31.1-1.1) ...
Selecting previously unselected package kubernetes-cni.
Preparing to unpack .../4-kubernetes-cni_1.5.1-1.1_amd64.deb ...
Unpacking kubernetes-cni (1.5.1-1.1) ...
Selecting previously unselected package kubelet.
Preparing to unpack .../5-kubelet_1.31.1-1.1_amd64.deb ...
Unpacking kubelet (1.31.1-1.1) ...
Setting up conntrack (1:1.4.8-1ubuntu1) ...
Setting up kubect1 (1.31.1-1.1) ...
Setting up cri-tools (1.31.1-1.1) ...
Setting up kubernetes-cni (1.5.1-1.1) ...
Setting up kubeadm (1.31.1-1.1) ...
Setting up kubelet (1.31.1-1.1) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

```

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.

```
ubuntu@ip-172-31-87-78:~$ sudo apt-mark hold kubelet kubeadm kubect1
```

kubelet set on hold.

kubeadm set on hold.

kubect1 set on hold.

```
ubuntu@ip-172-31-87-78:~$ sudo systemctl enable --now kubelet
```

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
```

```

ubuntu@ip-172-31-87-78:~$ sudo systemctl enable --now kubelet
ubuntu@ip-172-31-87-78:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
WARNING: [11:48:52.28251] - 4743 checks.go:1080] [preflight] WARNING: Couldn't create the interface used for talking to the container runtime: failed to create new CRI runtime service: validate service connection:
validate CRI v1 runtime API for endpoint "unix:///var/run/containerd/containerd.sock": rpc error: code = Unimplemented desc = unknown service runtime.v1.RuntimeService
[WARNING FileExisting-socat]: socat not found in system path
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
error execution phase preflight: [preflight] Some fatal errors occurred:
failed to create new CRI runtime service: validate service connection: validate CRI v1 runtime API for endpoint "unix:///var/run/containerd/containerd.sock": rpc error: code = Unimplemented desc = unknown service runtime.v1.RuntimeService[preflight] If you know what you are doing, you can make a check non-fatal with '--ignore-preflight-errors=...'
To see the stack trace of this error execute with --v5 or higher
ubuntu@ip-172-31-87-78:~$

```

```
sudo mkdir -p /etc/containerd
```

```
sudo containerd config default | sudo tee /etc/containerd/config.toml
```



```
to see the stack trace of this error execute with --v=3 or higher
ubuntu@ip-172-31-87-78:~$ sudo apt-get install -y containerd
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed:
  runc
The following packages will be REMOVED:
  containerd.io docker-ce
The following NEW packages will be installed:
  containerd runc
0 upgraded, 2 newly installed, 2 to remove and 139 not upgraded.
Need to get 47.2 MB of archives.
After this operation, 53.1 MB disk space will be freed.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 runc amd64 1.1.12-0ubuntu3.1 [8599 kB]
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 containerd amd64 1.7.12-0ubuntu4.1 [38.6 MB]
Fetched 47.2 MB in 1s (88.2 MB/s)
(Reading database ... 68064 files and directories currently installed.)
Removing docker-ce (5:27.3.1-1~ubuntu.24.04~noble) ...
Removing containerd.io (1.7.22-1) ...
Selecting previously unselected package runc.
(Reading database ... 68044 files and directories currently installed.)
Preparing to unpack .../runc_1.1.12-0ubuntu3.1_amd64.deb ...
Unpacking runc (1.1.12-0ubuntu3.1) ...
Selecting previously unselected package containerd.
Preparing to unpack .../containerd_1.7.12-0ubuntu4.1_amd64.deb ...
Unpacking containerd (1.7.12-0ubuntu4.1) ...
Setting up runc (1.1.12-0ubuntu3.1) ...
Setting up containerd (1.7.12-0ubuntu4.1) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-87-78:~$ █
```

```

ubuntu@ip-172-31-87-78:~$ sudo mkdir -p /etc/containerd
sudo containerd config default | sudo tee /etc/containerd/config.toml
ubuntu@ip-172-31-87-78:~$ sudo mkdir -p /etc/containerd
ubuntu@ip-172-31-87-78:~$ sudo containerd config default | sudo tee /etc/containerd/config.toml
disabled_plugins = []
imports = []
oom_score = 0
plugin_dir = ""
required_plugins = []
root = "/var/lib/containerd"
state = "/run/containerd"
temp = ""
version = 2

[cgroup]
  path = ""

[debug]
  address = ""
  format = ""
  gid = 0
  level = ""
  uid = 0

[grpc]
  address = "/run/containerd/containerd.sock"
  gid = 0
  max_recv_message_size = 16777216
  max_send_message_size = 16777216
  tcp_address = ""
  tcp_tls_ca = ""
  tcp_tls_cert = ""
  tcp_tls_key = ""
  uid = 0

```

sudo systemctl restart containerd

sudo systemctl enable containerd

sudo systemctl status containerd

```

ubuntu@ip-172-31-87-78:~$ sudo systemctl restart containerd
ubuntu@ip-172-31-87-78:~$ sudo systemctl enable containerd
ubuntu@ip-172-31-87-78:~$ sudo systemctl status containerd
* containerd.service - containerd container runtime
   Loaded: loaded (/usr/lib/systemd/system/containerd.service; enabled; preset: enabled)
   Active: active (running) since Sat 2024-09-21 12:48:02 UTC; 2min 6s ago
     Docs: https://containerd.io
   Main PID: 7830 (containerd)
      Tasks: 7
    Memory: 13.2M (peak: 14.0M)
       CPU: 295ms
    CGroup: /system.slice/containerd.service
            └─7830 /usr/bin/containerd

Sep 21 12:48:02 ip-172-31-87-78 containerd[7830]: time="2024-09-21T12:48:02.680681825Z" level=info msg="Start subscribing containerd event"
Sep 21 12:48:02 ip-172-31-87-78 containerd[7830]: time="2024-09-21T12:48:02.680713219Z" level=info msg="serving... address=/run/containerd/containerd.sock.ttrpc"
Sep 21 12:48:02 ip-172-31-87-78 containerd[7830]: time="2024-09-21T12:48:02.680720772Z" level=info msg="Start recovering state"
Sep 21 12:48:02 ip-172-31-87-78 containerd[7830]: time="2024-09-21T12:48:02.680745293Z" level=info msg="serving... address=/run/containerd/containerd.sock"
Sep 21 12:48:02 ip-172-31-87-78 containerd[7830]: time="2024-09-21T12:48:02.680770644Z" level=info msg="Start event monitor"
Sep 21 12:48:02 ip-172-31-87-78 containerd[7830]: time="2024-09-21T12:48:02.680779762Z" level=info msg="Start snapshots syncer"
Sep 21 12:48:02 ip-172-31-87-78 containerd[7830]: time="2024-09-21T12:48:02.680787308Z" level=info msg="Start cni network conf syncer for default"
Sep 21 12:48:02 ip-172-31-87-78 containerd[7830]: time="2024-09-21T12:48:02.680795250Z" level=info msg="Start streaming server"
Sep 21 12:48:02 ip-172-31-87-78 containerd[7830]: time="2024-09-21T12:48:02.680847770Z" level=info msg="containerd successfully booted in 0.027952s"
Sep 21 12:48:02 ip-172-31-87-78 systemd[1]: Started containerd.service - containerd container runtime.
ubuntu@ip-172-31-87-78:~$

```

sudo apt-get install -y socat

```

ubuntu@ip-172-31-87-78:~$ sudo apt-get install -y socat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libltdl7 libslirp0 pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  socat
0 upgraded, 1 newly installed, 0 to remove and 139 not upgraded.
Need to get 374 kB of archives.
After this operation, 1649 kB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 socat amd64 1.8.0-4build3 [374 kB]
Fetched 374 kB in 0s (13.2 MB/s)
Selecting previously unselected package socat.
(Reading database ... 68108 files and directories currently installed.)
Preparing to unpack .../socat_1.8.0-4build3_amd64.deb ...
Unpacking socat (1.8.0-4build3) ...
Setting up socat (1.8.0-4build3) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-87-78:~$

```

## Step 6: Initialize the Kubecluster

`sudo kubeadm init --pod-network-cidr=10.244.0.0/16`

```

ubuntu@ip-172-31-87-78:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W0921 12:58:06.057054 8434 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.8" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-87-78.kubernetes.kubernetes.default.kubernetes.default.svc.kubernetes.default.svc.cluster.local] and IPs [10.96.0.1 172.31.87.78]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-87-78.localhost] and IPs [172.31.87.78 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-87-78.localhost] and IPs [172.31.87.78 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Writing "admin.conf" kubeconfig file
[kubeconfig] Writing "super-admin.conf" kubeconfig file
[kubeconfig] Writing "kubelet.conf" kubeconfig file
[kubeconfig] Writing "controller-manager.conf" kubeconfig file
[kubeconfig] Writing "scheduler.conf" kubeconfig file
[etcd] Creating static Pod manifest for local etcd in "/etc/kubernetes/manifests"
[control-plane] Using manifest folder "/etc/kubernetes/manifests"
[control-plane] Creating static Pod manifest for "kube-apiserver"
[control-plane] Creating static Pod manifest for "kube-controller-manager"
[control-plane] Creating static Pod manifest for "kube-scheduler"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"

```

Copy the mkdir and chown commands from the top and execute them.

`mkdir -p $HOME/.kube`

`sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`

`sudo chown $(id -u):$(id -g) $HOME/.kube/config`

```
[mark-control-plane] Marking the node ip-172-31-87-78 as control-plane by adding the taints [node-role.kubernetes.io/control-plane:NoSchedule]
[bootstrap-token] Using token: ut4opm.bfe28j57nfncds8m
[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to get nodes
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] Configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy
```

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:  
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.87.78:6443 --token ut4opm.bfe28j57nfncds8m \
--discovery-token-ca-cert-hash sha256:be2108f5129c3aac922280e2219dcace14c20204be972d8823eacc59bf6528e6
ubuntu@ip-172-31-87-78:~$
```

```
--discovery-token-ca-cert-hash sha256:be2108f5129c3aac922280e2219dcace14c20204be972d8823eacc59bf6528e6
ubuntu@ip-172-31-87-78:~$ mkdir -p $HOME/.kube
ubuntu@ip-172-31-87-78:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
ubuntu@ip-172-31-87-78:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
ubuntu@ip-172-31-87-78:~$
```

Step 7: Now that the cluster is up and running, we can deploy our nginx server on this cluster. Apply this deployment file using this command to create a deployment  
**kubectl apply -f <https://k8s.io/examples/application/deployment.yaml>**

```
See kubectl apply --help for usage.
ubuntu@ip-172-31-87-78:~$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
ubuntu@ip-172-31-87-78:~$
```

```
daemonset.apps/kube-flannel-ds created
ubuntu@ip-172-31-87-78:~$ kubectl apply -f https://k8s.io/examples/application/deployment.yaml
deployment.apps/nginx-deployment created
ubuntu@ip-172-31-87-78:~$
```

**kubectl get pods**

```
deployment.apps/nginx-deployment created
ubuntu@ip-172-31-87-78:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-d556bf558-2nwj8    0/1     Pending   0           6m54s
nginx-deployment-d556bf558-vbnn6    0/1     Pending   0           6m54s
ubuntu@ip-172-31-87-78:~$
```

**POD\_NAME=\$(kubectl get pods -l app=nginx -o jsonpath="{.items[0].metadata.name}")**  
**kubectl port-forward \$POD\_NAME 8080:80**

Note : We have faced an error as pod status is pending so make it running run below commands

then again run above 2 commands.

**kubectl taint nodes --all node-role.kubernetes.io/control-plane-node/ip-172-31-20-171 untainted**

**kubectl get nodes**

Step 8: Verify your deployment

Open up a new terminal and ssh to your EC2 instance.

Then, use this curl command to check if the Nginx server is running.

```
ubuntu@ip-172-31-20-171:~$ curl --head http://127.0.0.1:8080
HTTP/1.1 200 OK
Server: nginx/1.14.2
Date: Sun, 15 Sep 2024 07:59:03 GMT
Content-Type: text/html
Content-Length: 612
Last-Modified: Tue, 04 Dec 2018 14:44:49 GMT
Connection: keep-alive
ETag: "5c0692e1-264"
Accept-Ranges: bytes

ubuntu@ip-172-31-20-171:~$
```

If the response is 200 OK and you can see the Nginx server name, your deployment was successful.

We have successfully deployed our Nginx server on our EC2 instance.

#### Conclusion:

In this experiment, we successfully installed Kubernetes on an EC2 instance and deployed an Nginx server using Kubectl commands. During the process, we encountered two main errors: the Kubernetes pod was initially in a pending state, which was resolved by removing the control-plane taint using `kubectl taint nodes --all`, and we also faced an issue with the missing containerd runtime, which was fixed by installing and starting containerd. We used a t2.medium EC2 instance with 2 CPUs to meet the necessary resource requirements for the Kubernetes setup and deployment.